

1.

To prove that $E_{agg} = 1/M E_{avg}$, under the given assumptions, let's start by expanding the expression for $E_{agg}(x)$:

$$E_{agg}(x) = E[(\{1/M \sum \epsilon_i(x)\}^2)]$$

Now, we can expand the square:

$$E_{agg}(x) = E[(\{1/M^2 \sum \epsilon_i(x) \sum \epsilon_j(x)\})]$$

Applying the expectation operator:

$$E_{agg}(x) = 1/M^2 \sum \sum E(\epsilon_i(x) \epsilon_j(x))$$

Now, let's apply the given assumptions:

Each of the errors has a 0 mean:

$$E(\epsilon_i(x)) = 0 \text{ for all } i$$

Errors are uncorrelated:

$$E(\epsilon_i(x) \epsilon_j(x)) = 0 \text{ for all } i \neq j$$

Using these assumptions, we can simplify the expression:

$$E_{agg}(x) = 1/M^2 \sum E(\epsilon_i(x)^2)$$

Now, we can rewrite the expression for E_{avg} using the given definition:

$$E_{avg} = 1/M \sum E(\epsilon_i(x)^2)$$

Comparing the expressions for $E_{agg}(x)$ and E_{avg} , we can see that:

$$E_{agg}(x) = 1/M E_{avg}$$

Thus, under the given assumptions of 0 mean errors and uncorrelated errors, the expected squared error of the aggregated model is equal to $1/M$ times the average expected squared error of the individual models.

2.

E_{agg} : Expected error of the aggregated model (a model that is a linear combination of M models)

E_{av} : The average expected error of individual models

$\epsilon_i(x)$: Error of the i -th model on a given input x

λ_i : The weight associated with the i -th model ($\lambda_i \geq 0$ and $\sum \lambda_i = 1$)

$E_{agg} = E(\sum \lambda_i \epsilon_i(x))$ and $E_{av} = M^{-1} \sum E(f(x_i))$. We want to show that $E(\sum \lambda_i \epsilon_i(x)) \leq M^{-1} \sum E(f(x_i))$.

Let's use Jensen's inequality, which states that for any convex function f , $f(\sum \lambda_i x_i) \leq \sum \lambda_i f(x_i)$, where x_i are real numbers and λ_i are non-negative real numbers such that $\sum \lambda_i = 1$.

$$f(\sum \lambda_i \epsilon_i(x)) \leq \sum \lambda_i f(\epsilon_i(x))$$

$$(\sum \lambda_i \epsilon_i(x))^2 \leq \sum \lambda_i (\epsilon_i(x))^2$$

$$E((\sum \lambda_i \epsilon_i(x))^2) \leq E(\sum \lambda_i (\epsilon_i(x))^2)$$

Since the left side represents the expected error of the aggregated model squared, it can be denoted as

$$E_{agg}^2.$$

On the right side, we can exchange the order of summation and expectation, as expectation is a linear operation:

$$E_{agg}^2 \leq \sum \lambda_i E((\epsilon_i(x))^2)$$

Now, note that $E((\epsilon_i(x))^2)$ is the expected squared error of the i -th model, which is related to the expected error of the i -th model.

Since $M^{-1} \sum E(f(x_i))$ represents the average expected error, and λ_i are the weights, we can rewrite the inequality as:

$$E_{agg}^2 \leq \sum \lambda_i E((\epsilon_i(x))^2)$$

We can rewrite $M^{-1} \sum E(f(x_i))$ as the weighted sum of expected squared errors of the individual models:

$$E_{av} = M^{-1} \sum E((\epsilon_i(x))^2)$$

Now, multiply both sides of the inequality by M :

$$ME_{agg}^2 \leq \sum \lambda_i E((e_i(x))^2)$$

Since $ME_{av} = \sum \lambda_i E((e_i(x))^2)$ and $M > 0$, we can compare the left sides of the two inequalities:

$$ME_{agg}^2 \leq ME_{av}$$

Divide both sides by M:

$$E_{agg}^2 \leq E_{av}$$

Since both E_{agg} and E_{av} are non-negative (errors are non-negative), we can take the square root of both sides:

$$\sqrt{E_{agg}^2} \leq \sqrt{E_{av}}$$

$$E_{agg} \leq E_{av}$$

3.

First, let's define the overall training error at the end of T steps, E:

$$E = (1/N) * \sum_i [H(x_i) \neq y_i]$$

where N is the total number of data points, x_i is the i th data point, y_i is the true label for x_i , and $H(x_i)$ is the final hypothesis for x_i .

Now, let's rewrite $H(x_i)$ using the equation provided:

$$H(x_i) = \text{sign}(\sum_{t=1}^{T} \alpha^t * h^t(x_i))$$

Since the weight for the point i at step $t+1$ is given by:

$$D_{t+1}(i) = D_t(i) / Z_t * \exp(-\alpha^t * h_t(i) * y(i))$$

We can compute the upper bound for the training error at each step as follows:

$$\sum_i [D_t(i) * [h_t(x_i) \neq y(x_i)]] \leq \exp(-2 * \gamma^2)$$

Next, let's express the overall training error E in terms of the weights D:

$$E = \sum_i [D_T(i) * [H(x_i) \neq y_i]] / N$$

Now, let's use the fact that the weights D are normalized at each step (i.e., $\sum_i D_t(i) = 1$):

$$E = \sum_i [D_T(i) * [H(x_i) \neq y_i]] / \sum_i D_T(i)$$

We can now upper bound the overall training error E by the sum of the upper bounds of the training error at each step:

$$E \leq \sum_i [D_t(i) * \exp(-2 * \gamma^2)] / \sum_i D_t(i)$$

Since $\sum_i D_t(i)$ is a normalization factor that ensures the weights sum up to 1, we can simplify the expression:

$$E \leq \exp(-2 * \sum_{t=1}^T \gamma^2)$$

Thus, we have proven that at the end of T steps, the overall training error E will be bounded by:

$$\exp(-2 * \sum_{t=1}^{T} \gamma^2)$$